

```
In [46]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [47]: df=pd.read_csv("https://gist.githubusercontent.com/stoicsapien1/72729ffbcd4b75b7bedd4b01b65bc9b5/raw/bf95a2e30fceb9f2ae990eac8379
```

```
In [48]: df.head()
```

```
Out[48]:
```

	Gender	Height	Weight
0	Male	73.847017	241.893563
1	Male	68.781904	162.310473
2	Male	74.110105	212.740856
3	Male	71.730978	220.042470
4	Male	69.881796	206.349801

```
In [49]: df["Gender"]=df["Gender"].map({"Male":0,"Female":1})
```

```
In [50]: df
```

Out[50]:


	Gender	Height	Weight
--	--------	--------	--------

0	0	73.847017	241.893563
1	0	68.781904	162.310473
2	0	74.110105	212.740856
3	0	71.730978	220.042470
4	0	69.881796	206.349801
...	...	...	...
9995	1	66.172652	136.777454
9996	1	67.067155	170.867906
9997	1	63.867992	128.475319
9998	1	69.034243	163.852461
9999	1	61.944246	113.649103

10000 rows × 3 columns

In [51]: `sns.heatmap(df.corr(),annot=True,cmap="coolwarm")`

Out[51]: <Axes: >



```
In [52]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()
```

```
In [53]: from sklearn.model_selection import train_test_split
```

```
In [54]: X=df.iloc[:,0:2]  
y=df.iloc[:, -1]
```

```
In [55]: X
```

Out[55]:

	Gender	Height
0	0	73.847017
1	0	68.781904
2	0	74.110105
3	0	71.730978
4	0	69.881796
...	...	...
9995	1	66.172652
9996	1	67.067155
9997	1	63.867992
9998	1	69.034243
9999	1	61.944246

10000 rows × 2 columns

In [56]:

```
y
```

Out[56]:

```
0    241.893563
1    162.310473
2    212.740856
3    220.042470
4    206.349801
```

```
...
```

```
9995   136.777454
9996   170.867906
9997   128.475319
9998   163.852461
9999   113.649103
```

Name: Weight, Length: 10000, dtype: float64

In [57]:

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=42)
```

In [58]:

```
lr.fit(X_train,y_train)
```

Out[58]:

LinearRegression ⓘ ?

LinearRegression()

In [59]:

X\_test

Out[59]:

	Gender	Height
6252	1	68.687232
4684	0	68.829334
1731	0	66.398128
4742	0	71.929340
4521	0	67.042903
...	...	...
4862	0	70.534667
7025	1	65.873106
7647	1	60.646988
7161	1	67.122422
73	0	67.595269

2500 rows × 2 columns

In [60]:

y\_test

```
Out[60]: 6252    173.115813
         4684    195.162042
         1731    170.217451
         4742    215.049660
         4521    182.721452
         ...
         4862    203.763292
         7025    155.915544
         7647    123.862430
         7161    165.324899
         73     163.108002
         Name: Weight, Length: 2500, dtype: float64
```

```
In [61]: y_pred=lr.predict(X_test)
```

```
In [62]: y_pred
```

```
Out[62]: array([165.57063635, 185.78881695, 171.26964188, ..., 117.55424148,
                156.22557983, 178.41897264])
```

```
In [63]: m=lr.coef_
```

```
In [64]: m
```

```
Out[64]: array([-19.36954901,  5.97200677])
```

```
In [68]: X_test["ypred"]=y_pred
```

```
In [69]: X_test
```

Out[69]:

	Gender	Height	ypred
6252	1	68.687232	165.570636
4684	0	68.829334	185.788817
1731	0	66.398128	171.269642
4742	0	71.929340	204.302076
4521	0	67.042903	175.120240
...	...	...	...
4862	0	70.534667	195.973080
7025	1	65.873106	148.764658
7647	1	60.646988	117.554241
7161	1	67.122422	156.225580
73	0	67.595269	178.418973

2500 rows × 3 columns

In [70]:

```
X_test["y_test"]=y_test
```

In [71]:

```
X_test
```

Out[71]:

	Gender	Height	ypred	y_test
6252	1	68.687232	165.570636	173.115813
4684	0	68.829334	185.788817	195.162042
1731	0	66.398128	171.269642	170.217451
4742	0	71.929340	204.302076	215.049660
4521	0	67.042903	175.120240	182.721452
...	...	...	...	...
4862	0	70.534667	195.973080	203.763292
7025	1	65.873106	148.764658	155.915544
7647	1	60.646988	117.554241	123.862430
7161	1	67.122422	156.225580	165.324899
73	0	67.595269	178.418973	163.108002

2500 rows × 4 columns

In [72]: `lr.predict([[0,74]])`

C:\Users\lucius seneca\anaconda3\lib\site-packages\sklearn\base.py:493: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names  
 warnings.warn(

Out[72]: `array([216.66807065])`

In [74]: `from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score`

In [75]: `r2_score(y_test,y_pred)`

Out[75]: `0.9026809109766419`

In [76]: `mean_absolute_error(y_test,y_pred)`

Out[76]: `7.993905209333191`




```
In [77]: mean_squared_error(y_test,y_pred)
```

```
Out[77]: 101.86825136044213
```

```
In [84]: plt.scatter(y_test,y_pred,marker="+",color="red")  
plt.xlabel("Actual")  
plt.ylabel("Predicted")  
plt.plot([0, 250], [0, 250], color='blue', linestyle='--', label='y=x')
```

```
Out[84]: [<matplotlib.lines.Line2D at 0x198c7f32050>]
```



```
In [ ]:
```